



The Vienna Computing Library

Florian Rudolf

Institute for Microelectronics
Institute for Analysis and Scientific Computing

Technische Universität Wien, Austria



CG Libs - Smart Libraries for Computer Graphics 2013


Talk Overview

Session 1


Introduction to ViennaCL
How-To: ViennaCL Basics

Session 2

How-To: Advanced ViennaCL
ViennaCL: Behind the curtain



Introduction to ViennaCL



What to expect

What is ViennaCL

OpenCL

History of ViennaCL

Goals of ViennaCL

Installation of ViennaCL




What Is ViennaCL?

What is it about the Name?

The beautiful city of **Vienna**

Open**CL**



History of OpenCL



OpenCL

Prior to 2008

OpenCL developed by Apple Inc.

2008

OpenCL working group formed at Khronos Group

OpenCL specification 1.0 released

2010

OpenCL 1.1 (multi-device, subbuffer manipulation)

2011

OpenCL 1.2 (device partitioning)




OpenCL Platform Model




Platform 1


OpenCL Platform Model



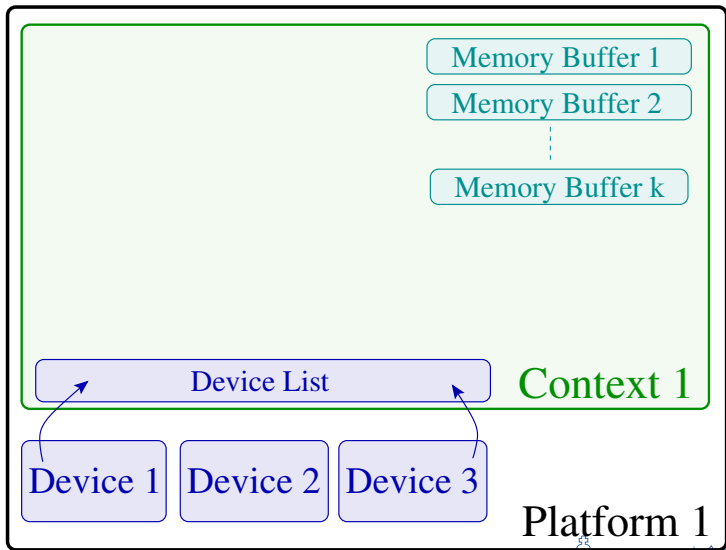
OpenCL Platform Model



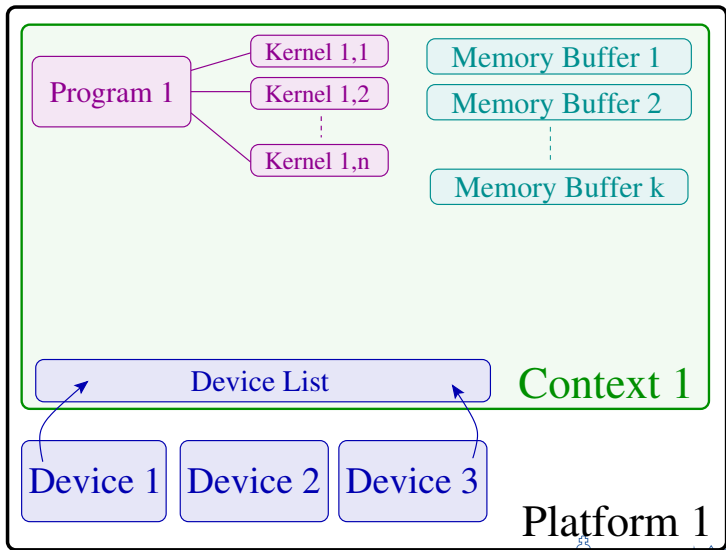
OpenCL Platform Model



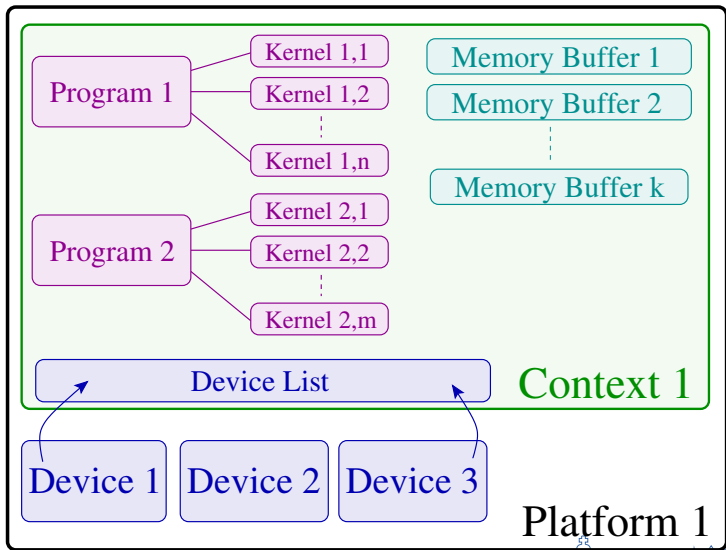
OpenCL Platform Model



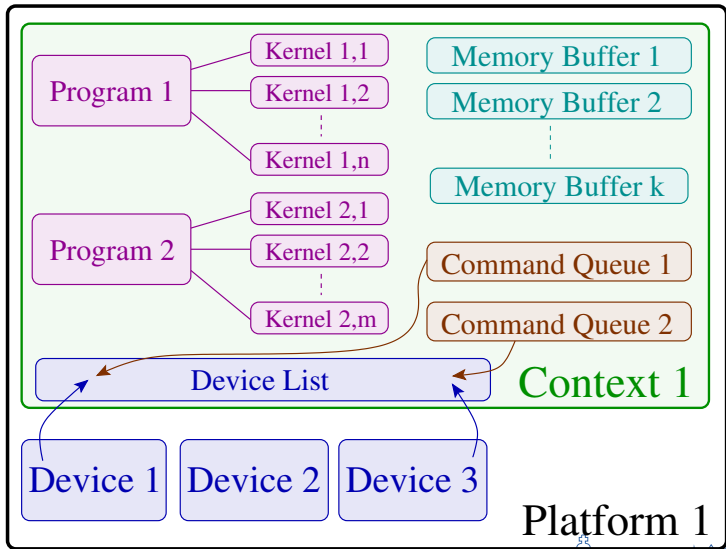
OpenCL Platform Model



OpenCL Platform Model



OpenCL Platform Model



OpenCL Host API

```
context = clCreateContextFromType(NULL, CL_DEVICE_TYPE_GPU, NULL, NULL,
    NULL);
queue = clCreateCommandQueue(context, NULL, 0, NULL);
memobjs[0] = clCreateBuffer(context, CL_MEM_READ_WRITE, sizeof(float)*2*
    num_entries, NULL, NULL);
memobjs[1] = clCreateBuffer(context, CL_MEM_READ_ONLY |
    CL_MEM_COPY_HOST_PTR, sizeof(float)*2*num_entries, srcA, NULL);

program = clCreateProgramWithSource(context, 1, &kernel_src, NULL, NULL);
clBuildProgram(program, 0, NULL, NULL, NULL, NULL);


kernel = clCreateKernel(program, "my_kernel", NULL);
clSetKernelArg(kernel, 0, sizeof(cl_mem), (void *)&memobjs[0]);
clSetKernelArg(kernel, 1, sizeof(cl_mem), (void *)&memobjs[1]);
clSetKernelArg(kernel, 2, sizeof(float)*(local_work_size[0]+1)*16, NULL);

global_work_size[0] = 128;
local_work_size[0] = 64;
clEnqueueNDRangeKernel(queue, kernel, 1, NULL, global_work_size,
    local_work_size, 0, NULL, NULL);
```

Issues

“Where is the error?”

Manage OpenCL handles



OpenCL Kernel Language

Sample Operation: Inplace Vector Addition

$$\begin{pmatrix} v_1^1 \\ v_1^2 \\ \vdots \\ v_1^n \end{pmatrix} += \begin{pmatrix} v_2^1 \\ v_2^2 \\ \vdots \\ v_2^n \end{pmatrix}$$

OpenCL Kernel

```
__kernel void inplace_add(  
    __global const float * vec1,  
    __global const float * vec2,  
    unsigned int size)  
{  
    for (unsigned int i = get_global_id(0);  
         i < size;  
         i += get_global_size(0))  
        vec1[i] += vec2[i];  
}
```


2010


April: Roots in the Master's Thesis of Florian Rudolf

May 28th: ViennaCL 1.0.0 released

November: 1000th download

December: ViennaCL 1.1.0

(BLAS level 3, refurbished OpenCL backend)



History of ViennaCL

2010

April: Roots in the Master's Thesis of Florian Rudolf

May 28th: ViennaCL 1.0.0 released

November: 1000th download

December: ViennaCL 1.1.0


(BLAS level 3, refurbished OpenCL backend)

2011

March: Accepted for Google Summer of Code

December: ViennaCL 1.2.0

(AMG, SPAI, FFT, QR, graph algorithms, structured matrices)



2012

March: Accepted for Google Summer of Code

May: Tutorial at NVIDIA GTC 2012

May: ViennaCL 1.3.0

(ranges and slices, Automated OpenCL kernels, eigen values, ILU0, SVD)

December: ViennaCL 1.4.0

(CUDA and host backend, initializer types)



History of ViennaCL

2012

March: Accepted for Google Summer of Code

May: Tutorial at NVIDIA GTC 2012

May: ViennaCL 1.3.0

(ranges and slices, Automated OpenCL kernels, eigen values, ILU0, SVD)

December: ViennaCL 1.4.0

(CUDA and host backend, initializer types)

2013

May: Accepted for Google Summer of Code


June: Tutorial at CGLibs



Goals of ViennaCL

About

High-level linear algebra C++ library
OpenMP, OpenCL, and CUDA backends
Header-only
Multi-platform




Dissemination

Free Open-Source MIT (X11) License
<http://viennacl.sourceforge.net/>
50-100 downloads per week

Design Rules

Reasonable default values
Compatible to Boost.uBLAS whenever possible
In doubt: clean design over performance



Core features

- Linear algebra, BLAS
- Solver (direct and iterative)
- Preconditioners

Additional features

- Fast Fourier Transform
- Eigenvalue computations
- QR factorization
- Bandwidth reduction
- Nonnegative matrix factorization



Interfaces to other libraries

Boost.uBLAS

Eigen

MTL4

Backends

CPU


OpenCL

CUDA

C++ library

Generic free functions

Expression templates



Installation of ViennaCL

Three Steps

Download from <http://viennacl.sourceforge.net/>

Unzip

Copy source folder

Dynamic Library?

ViennaCL is header-only


Linking depends on used backend (OpenMP, OpenCL, CUDA)

Sample Applications


22 tutorials

7 benchmarks

about 35 tests



How-To: ViennaCL Basics



What to expect

From BLAS to Boost.uBLAS to ViennaCL

Basic Types

OpenCL Kernels

Basic Usage: Data Management

Basic Usage: Algebra

Basic Usage: Solver

Summary



BLAS

Basic **L**inear **A**lgebra **S**ubprograms

De facto API standard

Low level interface

Boost.uBLAS

C++ API for BLAS

High level interface

Part of Boost libraries



From BLAS to Boost.uBLAS to ViennaCL

Consider Existing CPU Code (Boost.uBLAS)

```
using namespace boost::numeric::ublas;

matrix<double> A(1000, 1000);
vector<double> x(1000), rhs(1000);


/* Fill A, x, rhs here */

rhs += 2.0 * x;
double val = inner_prod(x, rhs);
matrix += val * outer_prod(x, rhs);

x = solve(A, rhs, upper_tag()); // Upper triangular solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```

High-Level Code with Syntactic Sugar



Porting the previous code to GPU

“How much time will you need?”

1 minute?

1 hour?

1 day?


Quality Criteria

Working hours spent

Performance

Correctness

High-level description (Maintainability)



Consider Existing CPU Code (Boost.uBLAS)

```
using namespace boost::numeric::ublas;

matrix<double> A(1000, 1000);
vector<double> x(1000), rhs(1000);

/* Fill A, x, rhs here */

rhs += 2.0 * x;
double val = inner_prod(x, rhs);
matrix += val * outer_prod(x, rhs);

x = solve(A, rhs, upper_tag()); // Upper triangular solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```



Previous Code Snippet Rewritten with ViennaCL

```
using namespace viennacl;
using namespace viennacl::linalg;

matrix<double> A(1000, 1000);
vector<double> x(1000), rhs(1000);

/* Fill A, x, rhs here */

rhs += 2.0 * x;
double val = inner_prod(x, rhs);
matrix += val * outer_prod(x, rhs);

x = solve(A, rhs, upper_tag()); // Upper triangular solver

std::cout << " 2-norm: " << norm_2(x) << std::endl;
std::cout << "sup-norm: " << norm_inf(x) << std::endl;
```



ViennaCL in addition provides iterative solvers

```
using namespace viennacl;  
using namespace viennacl::linalg;  
  
compressed_matrix<double> A(1000, 1000);  
vector<double> x(1000), rhs(1000);  
  
/* Fill A, x, rhs here */  
  
x = solve(A, rhs, cg_tag()); // Conjugate Gradient solver  
x = solve(A, rhs, bicgstab_tag()); // BiCGStab solver  
x = solve(A, rhs, gmres_tag()); // GMRES solver
```

No iterative solvers available in uBLAS...



Thanks to interface compatibility

```
using namespace boost::numeric::ublas;  
using namespace viennacl::linalg;  
  
compressed_matrix<double> A(1000, 1000);  
vector<double> x(1000), rhs(1000);  
  
/* Fill A, x, rhs here */  
  
x = solve(A, rhs, cg_tag()); // Conjugate Gradient solver  
x = solve(A, rhs, bicgstab_tag()); // BiCGStab solver  
x = solve(A, rhs, gmres_tag()); // GMRES solver
```



OpenCL kernels

OpenCL kernels are needed for actual computation

Provided by ViennaCL

Support for expression templates

Automatic kernel generation

Each of the following commands launches a separate OpenCL kernel

```
v1 = v2;  
v1 += v2;  
v1 -= v2;  
v1 = alpha * v2;  
v1 += alpha * v2;  
v1 -= alpha * v2;  
v1 *= alpha;  
v1 /= alpha;
```



OpenCL kernels have to be compiled at run time

- OpenCL JIT compiler

- Kernels can be grouped in programs

Compilation strategies

- Each kernel individually: several milliseconds per kernel

- All at once: Takes seconds

- In groups: Compile groups of kernels whenever potentially needed



OpenCL kernels

Consider the following code

```
int main(){
    vector<double> x(100), y(100);
    matrix<double> A(100, 100), B(100, 100);
    matrix<double, column_major> C(100, 100);

    x += 3.1415 * y;
    C = prod(trans(A), B);
    y = prod(C, x);


    std::cout << y << std::endl;
}
```

OpenCL Kernels

Which kernels are compiled?

When are they compiled?

Hint: Program compiles and executes normally



OpenCL kernels

Consider the following code


```
int main() {  
    vector<double> x(100), y(100);  
    matrix<double> A(100, 100), B(100, 100);  
    matrix<double, column_major> C(100, 100);  
  
    x += 3.1415 * y;  
    C = prod(trans(A), B);  
    y = prod(C, x);  
  
    std::cout << y << std::endl;  
}
```

OpenCL Kernels

Which kernels are compiled?

When are they compiled?

Hint: Program compiles and executes normally



OpenCL kernels

Consider the following code


```
int main() {  
    vector<double> x(100), y(100);  
    matrix<double> A(100, 100), B(100, 100);  
    matrix<double, column_major> C(100, 100);  
  
    x += 3.1415 * y;  
    C = prod(trans(A), B);  
    y = prod(C, x);  
  
    std::cout << y << std::endl;  
}
```

OpenCL Kernels

Which kernels are compiled?

When are they compiled?

Hint: Program compiles and executes normally



Special case: Matrix-Matrix product

Result matrix: Row/Column major memory layout

First factor: Row/Column major, possibly transposed

Second factor: Row/Column major, possibly transposed

Leads to 32 different kernels for matrix-matrix multiplication

Compiled separately on request



Basic Types

Supported types

Scalar

Vector

Dense matrix

Sparse matrix

Structured matrix

Numeric Types

float

double



Scalars


Represents a single scalar value on the computing device

Behave like underlying type

Implicit cast to underlying type

Potentially expensive (Overhead!)

```
viennacl::scalar<NumericType> gpu_scalar;  
viennacl::scalar<float> gpu_float;  
viennacl::scalar<double> gpu_double;
```



Basic Types

Scalars

```
float cpu_float = 42.0f;
double cpu_double = 13.7603;
viennacl::scalar<float> gpu_float(3.1415f);
viennacl::scalar<double> gpu_double = 2.71828;

// conversions and t
cpu_float = gpu_float;
// automatic transfer and conversion
gpu_float = cpu_double;

cpu_float = gpu_float * 2.0f;
cpu_double = gpu_float - cpu_float;
```



Vectors

Represents a vector on the computing device

Operator overloading

Alignment support

```
viennacl::vector<NumericType> gpu_vector;  
  
viennacl::vector<float> gpu_float_vector;  
viennacl::vector<double> gpu_double_vector;
```



Basic Types

Vectors

```
std::vector<ScalarType> stl_vec(10);
viennacl::vector<ScalarType> vcl_vec(10);

//fill the STL vector:
for (unsigned int i=0; i<vector_size; ++i)
    stl_vec[i] = i;

//copy content to GPU vector (recommended initialization)
copy(stl_vec.begin(), stl_vec.end(), vcl_vec.begin());

//manipulate GPU vector here
vcl_vec *= 4.2;

//copy content from GPU vector back to STL vector
copy(vcl_vec.begin(), vcl_vec.end(), stl_vec.begin());
```



Basic Types

Vectors

```
std::vector<ScalarType> stl_vec(10);  
viennacl::vector<ScalarType> vcl_vec(10);  
  
//fill the STL vector:  
for (unsigned int i=0; i<vector_size; ++i)  
    stl_vec[i] = i;  
  
//copy content to GPU vector (recommended initialization)  
copy(stl_vec, vcl_vec);  
  
//manipulate GPU vector here  
vcl_vec *= 4.2;  
  
//copy content from GPU vector back to STL vector  
copy(vcl_vec, stl_vec);
```




Alignment

Default = 1

Template parameter

In ViennaCL 1.5.0 deprecated (will be runtime parameter)

```
viennacl::vector<NumericType, Alignment>  
    gpu_vector_with_alignment;  
  
viennacl::vector<float, 4>  
    gpu_float_vector_with_alignment;  
  
viennacl::vector<double, 8>  
    gpu_double_vector_with_alignment;
```



Vector initializer

$$\text{unit_vector} \Rightarrow \mathbf{u}_{N,i} = \begin{cases} 1 & \text{if } i = N; \\ 0 & \text{else} \end{cases}$$

$$\text{zero_vector} \Rightarrow \mathbf{z}_i = 0$$

$$\text{scalar_vector} \Rightarrow \mathbf{s}_i = s$$

```
viennacl::unit_vector<NumericType>(size, index);  
  
viennacl::zero_vector<NumericType>(size);  
  
viennacl::scalar_vector<NumericType>(size, scalar);
```



Vector initializer

```
viennacl::vector<ScalarType> vcl_vec =  
    viennacl::unit_vector<ScalarType>(10, 5);  
// Creates the vector (0, 0, 0, 0, 1, 0, 0, 0, 0, 0)  
  
viennacl::vector<ScalarType> vcl_vec =  
    viennacl::zero_vector<ScalarType>(10);  
// Creates the vector (0, 0, 0, 0, 0, 0, 0, 0, 0, 0)  
  
viennacl::vector<ScalarType> vcl_vec =  
    viennacl::scalar_vector<ScalarType>(6, 1.5);  
// Creates the vector (1.5, 1.5, 1.5, 1.5, 1.5, 1.5)
```



Basic Types

Dense matrix

Represents a dense matrix on the computing device

Dense matrix \Rightarrow zero elements are rare

Alignment support (same as vectors)

Row major or column major

```
viennacl::matrix<NumericType> gpu_matrix;  
viennacl::matrix<NumericType, Scheme, Alignment>  
    gpu_matrix_with_scheme_and_alignment;  
  
viennacl::matrix<float> gpu_float_matrix;  
viennacl::matrix<float, row_major, 4>  
    row_major_gpu_float_matrix_with_alignment;  
  
viennacl::matrix<double> gpu_double_matrix;  
viennacl::matrix<double, column_major, 8>  
    column_major_gpu_double_matrix_with_alignment;
```

Dense matrix

```
//set up a 3 by 5 matrix:  
viennacl::matrix<float> vcl_matrix(4, 5);  
//fill it up:  
vcl_matrix(0,2) = 1.0  
vcl_matrix(1,2) = -1.5;  
vcl_matrix(2,0) = 4.2;  
vcl_matrix(3,4) = 3.1415;
```



Matrix initializer

$$\text{identity_matrix} \Rightarrow \mathbf{I}_{i,j} = \begin{cases} d & \text{if } i = j; \\ 0 & \text{else} \end{cases}$$

$$\text{zero_matrix} \Rightarrow \mathbf{Z}_{i,j} = 0$$

$$\text{scalar_matrix} \Rightarrow \mathbf{S}_{i,j} = s$$


```
viennacl::identity_matrix<NumericType>(size, diagonal);  
  
viennacl::zero_matrix<NumericType>(row_number,  
    column_number);  
  
viennacl::scalar_matrix<NumericType>(row_number,  
    column_number, scalar);
```



Matrix initializer

```
viennacl::matrix<ScalarType> vcl_mat =  
    viennacl::identity_matrix<ScalarType>(4, 1);
```


Creates the following matrix:

$$\begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$


Matrix initializer

```
viennacl::matrix<ScalarType> vcl_mat =  
    viennacl::zero_matrix<ScalarType>(3, 5);
```


Creates the following matrix: $\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$



Matrix initializer

```
viennacl::matrix<ScalarType> vcl_mat =  
    viennacl::scalar_matrix<ScalarType>(4, 3, 4.2);
```

Creates the following matrix:

$$\begin{pmatrix} 4.2 & 4.2 & 4.2 \\ 4.2 & 4.2 & 4.2 \\ 4.2 & 4.2 & 4.2 \\ 4.2 & 4.2 & 4.2 \end{pmatrix}$$


Sparse matrix

Represents a sparse matrix on the computing device

Sparse matrix \Rightarrow zero elements are frequent

Alignment support (same as vectors and dense matrices)

Sparse matrix types

Coordinate matrix

Compressed matrix

ELL matrix


Hybrid matrix



Basic Types

Coordinate matrix


Easy to setup/extend



Compressed matrix

Less memory required

Fast matrix-vector multiplication



ELL matrix

- Similar to compressed matrix

- Fixed number of non-zero values per row

- No row jumper array required

Hybrid matrix

- Combination of compressed matrix and ELL matrix



Compressed matrix

```
//set up a sparse 4 by 5 matrix on the CPU:
std::vector< std::map< unsigned int, float> >
    cpu_sparse_matrix(4);

//fill it up:
cpu_sparse_matrix[0][2] = 1.0;
cpu_sparse_matrix[1][2] = -1.5;
cpu_sparse_matrix[3][0] = 4.2;

//set up a sparse ViennaCL matrix:
viennacl::compressed_matrix<float> sparse_matrix(4, 5);

//copy to OpenCL device:
copy(cpu_sparse_matrix, sparse_matrix);

//copy back to CPU:
copy(sparse_matrix, cpu_sparse_matrix);
```

Structured matrix

Dense matrices but with special structure

Access to one element might change other elements

Supported types

Circulant matrix

Hankel matrix

Toeplitz matrix

Vandermonde matrix



Structured matrix

Circulant matrix

Hankel matrix

Toeplitz matrix

Vandermonde matrix

$$\begin{pmatrix} c_0 & c_{n-1} & \dots & c_2 & c_1 \\ c_1 & c_0 & c_{n-1} & & c_2 \\ \vdots & c_1 & c_0 & \ddots & \vdots \\ c_{n-2} & & \ddots & \ddots & c_{n-1} \\ c_{n-1} & c_{n-2} & \dots & c_1 & c_0 \end{pmatrix}$$



Structured matrix


Circulant matrix

Hankel matrix

Toeplitz matrix

Vandermonde matrix

$$\begin{pmatrix} a & b & c & d \\ b & c & d & e \\ c & d & e & f \\ d & e & f & g \end{pmatrix}$$



Structured matrix

Circulant matrix

Hankel matrix

Toeplitz matrix

Vandermonde matrix

$$\begin{pmatrix} a & b & c & d \\ e & a & b & c \\ f & e & a & b \\ g & f & e & a \end{pmatrix}$$



Basic Types

Structured matrix

Circulant matrix

Hankel matrix

Toeplitz matrix

Vandermonde matrix

$$\begin{pmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \vdots & \vdots & \vdots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{pmatrix}$$



Basic Usage: Data Management

How to access/transfer ViennaCL vectors/matrices elements?

Direct element access

```
vector<ScalarType> vcl(10);  
  
for (int i = 0; i < 10; ++i)  
    vcl(i) = i;  
  
for (int i = 0; i < 10; ++i)  
    std::cout << vcl(i) << std::endl;
```



Basic Usage: Data Management

How to access/transfer ViennaCL vectors/matrices elements?

Direct element access

Iterator

```
vector<ScalarType> vcl(10);

ScalarType tmp = 0;
for (vector<ScalarType>::iterator it = vcl.begin();
      it != vcl.end(); ++it, tmp += 42.0)
    *it = tmp;

for (vector<ScalarType>::iterator it = vcl.begin();
      it != vcl.end(); ++it)
    std::cout << *it < std::endl;
```



Basic Usage: Data Management

How to access/transfer ViennaCL vectors/matrices elements?

Direct element access

Iterator

Copy functions

```
std::vector<ScalarType> cpu(10);  
viennacl::vector<ScalarType> vcl(10);  
  
for (int i = 0; i < 10; ++i)  
    cpu[i] = i;  
  
viennacl::copy( cpu.begin(), cpu.end(), vcl.begin() );  
  
viennacl::copy( vcl.begin(), vcl.end(), cpu.begin() );
```



Basic Usage: Data Management

How to access/transfer ViennaCL vectors/matrices elements?

Direct element access

Iterator

Copy functions

```
std::vector<ScalarType> cpu(10);  
viennacl::vector<ScalarType> vcl(10);  
  
for (int i = 0; i < 10; ++i)  
    cpu[i] = i;  
  
viennacl::copy( cpu, vcl );  
  
viennacl::copy( vcl, cpu );
```



Granularity of Operations

Filling a vector with data

```
viennacl::vector<double> v(10000);  
  
for (size_t i=0; i<v.size(); ++i)  
    v(i) = i;
```



Basic Usage: Data Management

Granularity of Operations

Filling a vector with data

```
viennacl::vector<double> v(10000);  
  
for (size_t i=0; i<v.size(); ++i)  
    v(i) = i;
```

GPU Computing Is Fast, Right?

Execution time: 1 sec (approx)

std::vector: < 1 ms



Basic Usage: Data Management

Granularity of Operations

Transfer is done for each element separately

High overhead, similar to scalar operations

How to avoid the pitfall

Use temporary vector

Use copy functions

```
viennacl::vector<double> v(10000);  
std::vector<double> cpu_v( v.size() );  
  
for (size_t i=0; i<cpu_v.size(); ++i)  
    cpu_v(i) = i;  
  
viennacl::copy(cpu_v, v);
```

Basic Usage: Data Management

How to access/transfer ViennaCL vectors/matrices elements?

Direct element access \Rightarrow Bad idea!

Iterator \Rightarrow Bad idea!

Copy functions \Rightarrow **Good idea!**

```
std::vector<ScalarType> cpu_vec(10);  
viennacl::vector<ScalarType> vcl_vec(10);  
  
for (int i = 0; i < 10; ++i)  
    cpu_vec[i] = i;  
  
viennacl::copy( cpu_vec, vcl_vec );  
  
viennacl::copy( vcl_vec, cpu_vec );
```



Fast copy

copy does not require linear arrays \Rightarrow temporary required

```
std::list<ScalarType> cpu_vec(10);  
viennacl::vector<ScalarType> vcl_vec(10);  
  
for (int i = 0; i < 10; ++i)  
    cpu_vec[i] = i;  
  
viennacl::copy( cpu_vec, vcl_vec );  
  
viennacl::copy( vcl_vec, cpu_vec );
```



Fast copy

copy does not require linear arrays \Rightarrow temporary required

If container is linear memory \Rightarrow use fast_copy instead

```
std::vector<ScalarType> cpu_vec(10);  
viennacl::vector<ScalarType> vcl_vec(10);  
  
for (int i = 0; i < 10; ++i)  
    cpu_vec[i] = i;  
  
viennacl::fast_copy( cpu_vec, vcl_vec );  
  
viennacl::fast_copy( vcl_vec, cpu_vec );
```



Algebra operations

Trivial operations done by **operator overloading**

Scalar multiplication support for vector and matrices

Inner product and norm support for vectors


Matrix transpose using function **trans**

Matrix-vector and matrix-matrix multiplication using function **prod**



Scalar operations

```
NumericType s1, s2;  
viennacl::scalar<NumericType> vcl_s1, vcl_s2, vcl_s3;  
  
vcl_s1  = 5;  
vcl_s2  = vcl_s1 * 3;  
vcl_s3 -= vcl_s1 + (vcl_s2 / 4);  
  
s1      = vcl_s3;  
s2      = vcl_s1 - vcl_s2 * 2;
```



Vector operations

```
viennacl::scalar<NumericType> vcl_s1, vcl_s2;  
viennacl::vector<NumericType> v1(10), v2(10), v3(10);  
  
v2  = vcl_s1 * v1 + v2  
v3 += vcl_s1 * v2;  
  
v3  = 0.5 * v2 - v3;
```



Vector operations

```
NumericType s1, s2;  
viennacl::scalar<NumericType> vcl_s1, vcl_s2;  
viennacl::vector<NumericType> v1(10), v2(10), v3(10);  
  
vcl_s1 = viennacl::linalg::inner_prod(v1, v2);  
s1      = viennacl::linalg::inner_prod(v1, v2);  
  
s1      = viennacl::linalg::norm_1(v1);  
vcl_s2 = viennacl::linalg::norm_2(v2);  
s2      = viennacl::linalg::norm_inf(v3);
```



Matrix operations

```
viennacl::scalar<NumericType> vcl_s1, vcl_s2;  
viennacl::matrix<NumericType> M1(10, 10),  
    M2(10, 10), M3(10,10);  
  
M2  = vcl_s1 * M1 + M2  
M3 += vcl_s2 * M2;  
  
M3  = 0.5 * M2 - M3;  
  
M3  = viennacl::trans(M2); // Transposed matrix  
  
// Matrix-matrix product  
M1  = viennacl::linalg::prod( M2, M3 );  
M1  = viennacl::linalg::prod( M2, viennacl::trans(M3) );
```



GEMM: ViennaCL vs. CUBLAS

```
// ViennaCL
M1 = vcl_s1 * prod( M2, trans(M3) ) + vcl_s2 * M3;

// CUBLAS
cublasStatus_t cublasDgemm(handle,
    transa, transb,
    m, n, k,
    alpha,
    A, lda,
    B, ldb,
    beta,
    C, ldc);
```



Matrix-vector operations

```
viennacl::vector<NumericType> v1(10), v2(10), v3(20);  
viennacl::matrix<NumericType> M1(10, 10), M2(20, 10);  
  
v1 = viennacl::linalg::prod( M1, v2 );  
v1 = viennacl::linalg::prod( viennagrid::trans(M1), v2 );  
v3 = viennacl::linalg::prod( M2, v2 );  
v1 = viennacl::linalg::prod( M1, v3 );  
    // ERROR! dimension mismatch
```



Solving a system of linear equations

$$Ax = b$$

A and b given, x is unknown

Common problem in mathematics

Types of solver

Direct solver

Iterative solver



Direct solver

Solving the system directly

e.g. Gaussian elimination with pivoting

Iterative solver


Solving using an iterative process

Convergence relies on matrix properties

Recommended for large and sparse systems

No write operation needed on matrix

(most only use matrix-vector multiplication)



Basic Usage: Solver

Direct solver

LU factorization

No pivoting (work in progress)

```
using namespace viennacl::linalg;

viennacl::matrix<float> vcl_matrix;
viennacl::vector<float> vcl_rhs, vcl_result;
/* Set up matrix and vectors here */

//solution of an upper triangular system:
vcl_result = solve(vcl_matrix, vcl_rhs, upper_tag());
//solution of a lower triangular system:
vcl_result = solve(vcl_matrix, vcl_rhs, lower_tag());

//solution of a full system right into the vector vcl_rhs:
lu_factorize(vcl_matrix);
lu_substitute(vcl_matrix, vcl_rhs);
```


Iterative solver

Conjugate Gradient (CG)

Stabilized Bi-CG (BiCGStab)

Generalized Minimum Residual (GMRES)

Method		Matrix class	ViennaCL
Conjugate Gradient (CG)		symmetric positive definite	<code>y = solve(A, x, cg_tag());</code>
Stabilized Bi-CG (BiCGStab)		non-symmetric	<code>y = solve(A, x, bicgstab_tag());</code>
Generalized Minimum Residual (GMRES)		general	<code>y = solve(A, x, gmres_tag());</code>



Basic Usage: Solver

Iterative solver

Solver configuration via tag

```
using namespace viennacl::linalg;

// conjugate gradient solver with tolerance 1e10
// and at most 100 iterations:
viennacl::linalg::cg_tag custom_cg(1e-10, 100);

vcl_result = solve(vcl_matrix, vcl_rhs, custom_cg);

//print number of iterations taken and estimated error:
cout << "No. of iters: " << custom_cg.iters() << endl;
cout << "Est. error: " << custom_cg.error() << endl;
```



What have we learned?

ViennaCL provides interface compatibility with Boost.uBLAS

Basic types of ViennaCL

How OpenCL kernels are used

How to transfer data to and from ViennaCL


How to work with ViennaCL types

Simple algebraic operations

Solver for systems of linear equations



How-To: Advanced ViennaCL



What to expect

- Subvectors and Submatrices
- Escaping the Curse of Temporaries
- Interface: Eigen
- Performance
- Summary



Subvectors and Submatrices

Important for Many Algorithms

LU, Cholesky

QR, SVD

Two sub-types available

Range $[a, b)$

Slice $a:\text{inc}:\text{size}$

Ranges and slices are proxies

Read- and writeable



Subvectors and Submatrices

Range example

```
std::size_t lower_bound = 1;
std::size_t upper_bound = 7;
viennacl::range r(lower_bound, upper_bound);

typedef viennacl::vector<ScalarType> VectorType;
typedef viennacl::matrix<ScalarType> MatrixType;

// v[1:6]
viennacl::vector_range<VCLVectorType> v_sub(v, r);
// M[1:6,1:6]
viennacl::matrix_range<VCLMatrixType> M_sub(M, r, r);
```




Subvectors and Submatrices

Slice example

```
std::size_t start = 2;
std::size_t stride = 3;
std::size_t size = 5
viennacl::slice s(start, stride, size);

typedef viennacl::vector<ScalarType> VectorType;
typedef viennacl::matrix<ScalarType> MatrixType;

// v[2, 5, 8, 11, 14]
viennacl::vector_slice<VCLVectorType> v_sub(v, r);
// M[2,2], ..., M[2,14], ..., M[14,2], ..., M[14,14]
viennacl::matrix_slice<VCLMatrixType> M_sub(M, r, r);
```




Convenience Functions

```
viennacl::vector<ScalarType> v1(4), v2(2);  
viennacl::matrix<ScalarType> M1(4,4), M2(2,2);  
  
range r(0, 2);  
slice s(0, 2, 2);  
  
v2 = project(v1, r);  
project(v1, s) = v2;  
  
M2 = project(M1, r, r);  
viennacl::copy(M2, project(M1, s, s) );
```




Subvectors and Submatrices

Copy Headaches: Row-Major



Copy Headaches: Column-Major



Escaping the Curse of Temporaries

Simple BLAS Level 1 Operation

Consider

```
vec1 = vec2 + alpha * vec3 - beta * vec4;
```

With naive C++, this is equivalent to

```
tmp1 <- alpha * vec3  
tmp2 <- beta * vec4;  
tmp3 <- tmp1 - tmp2;  
tmp4 <- vec2 + tmp3;  
vec1 <- tmp4;
```

Temporaries Lead to Poor Performance

Costly on CPUs, extremely expensive on GPUs

Use expression templates for avoiding temporaries



Escaping the Curse of Temporaries

Vector Addition


```
x = y + z;
```

Naive Operator Overloading

```
vector<T> operator+(vector<T> & v, vector<T> & w);
```

$t \leftarrow y + z, x \leftarrow t$

Temporaries are extremely expensive!



Escaping the Curse of Temporaries

Expression Templates

```
vector_expr<vector<T>, op_plus, vector<T> >  
operator+(vector<T> & v, vector<T> & w) { ... }  
  
vector::operator=(vector_expr<...> const & e) {  
    viennacl::linalg::avbv(*this, 1,e.lhs(), 1,e.rhs());  
}
```

Allows to Avoid a Significant Amount of Temporaries

Covers most frequent cases

Influence on compilation times moderate




Escaping the Curse of Temporaries

Expression templates have their limitations

```
viennacl::vector<NumericType> v;  
viennacl::matrix<NumericType> M;  
  
v = viennacl::linalg::prod(M, v);
```

Temporary object is required here!

ViennaCL detects such cases and takes care of it



Data transfer

Like transfer from and to std container

Using provided copy functions

Interface compatibility

ViennaCL algorithms work with Eigen

e.g.: iterative solver



Data transfer: vectors

```
#define VIENNACL_HAVE_EIGEN

Eigen::VectorXd eigen_vector(dim);

// fill Eigen objects in a very sophisticated way with
numbers here

viennacl::vector<double> viennacl_vector(dim);

// copy data from Eigen objects to ViennaCL objects
viennacl::copy(eigen_vector, viennacl_vector);

// do some heavy linear algebra with ViennaCL here

// copy back to Eigen:
viennacl::copy(viennacl_vector, eigen_vector);
```



Data transfer: dense matrix

```
#define VIENNACL_HAVE_EIGEN

Eigen::MatrixXd eigen_densematrix(dim, dim);

// fill Eigen objects in a very sophisticated way with  
  numbers here

viennacl::matrix<double> viennacl_densematrix(dim, dim);

// copy data from Eigen objects to ViennaCL objects
viennacl::copy(eigen_densematrix, viennacl_densematrix);

// do some heavy linear algebra with ViennaCL here

// copy back to Eigen:
viennacl::copy(viennacl_densematrix, eigen_densematrix);
```



Interface: Eigen

Data transfer: sparse matrix

```
#define VIENNACL_HAVE_EIGEN

Eigen::SparseMatrix<double, Eigen::RowMajor>
    eigen_sparsematrix(dim, dim);

// fill Eigen objects in a very sophisticated way with
numbers here

viennacl::compressed_matrix<double> viennacl_sparsematrix(
    dim, dim);

// copy data from Eigen objects to ViennaCL objects
viennacl::copy(eigen_sparsematrix, viennacl_sparsematrix);

// do some heavy linear algebra with ViennaCL here

// copy back to Eigen:
viennacl::copy(viennacl_sparsematrix, eigen_sparsematrix);
```

Interface: Eigen

Interface Compatibility: iterative solver

```
#define VIENNACL_HAVE_EIGEN
using namespace viennacl::linalg;

Eigen::SparseMatrix<double, Eigen::RowMajor>
    matrix(dim, dim);
Eigen::VectorXd rhs(dim);
Eigen::VectorXd result(dim);
// fill eigen_matrix and eigen_rhs here

// Solve system using CG from ViennaCL
result = solve(matrix, rhs, cg_tag());

// Solve system using BiCGStab from ViennaCL
result = solve(matrix, rhs, bicgstab_tag());

// Solve system using GMRES from ViennaCL
result = solve(matrix, rhs, gmres_tag());
```

Granularity of Operations

Solving linear systems

```
viennacl::matrix<double> mat(N, N);  
viennacl::vector<double> rhs(N);  
  
for (size_t i=0; i<1000; ++i)  
{  
    viennacl::vector<double> result  
        = solve(mat, rhs, bicgstab_tag());  
    /* process result */  
}
```

Why Is There No Speed-Up?



Granularity of Operations

Solving linear systems

```
viennacl::matrix<double> mat(N, N);  
viennacl::vector<double> rhs(N);  
  
for (size_t i=0; i<1000; ++i)  
{  
    viennacl::vector<double> result  
        = solve(mat, rhs, bicgstab_tag());  
    /* process result */  
}
```

Why Is There No Speed-Up?

$$N = 3$$



Performance


Lets take a look



Performance


Sample Operation

$$v_1 \leftarrow v_2$$




OpenCL Kernel Launch Overhead

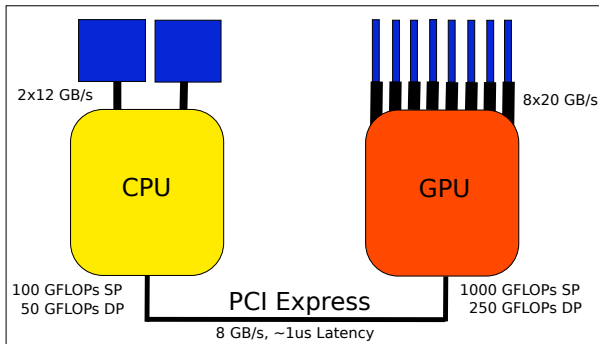
10 – 100 μ s



GPUs: Disillusion - Computing Architecture Schematic



GPUs: Disillusion - Computing Architecture Schematic




Good for large FLOP-intensive tasks, high memory bandwidth
PCI-Express can be a bottleneck

» 10-fold speedups (usually) not backed by hardware


Performance

Some benchmarks - vector addition




Performance

Some benchmarks - vector addition




Performance

Some benchmarks - vector addition




Performance

Some benchmarks - vector addition




Performance

Some benchmarks - vector addition




Performance

Some benchmarks - vector addition




Performance

Some benchmarks - vector addition




Performance

Some benchmarks - CG solver




Performance

Some benchmarks - CG solver




Performance

Some benchmarks - CG solver




Performance

Some benchmarks - CG solver




Performance

Some benchmarks - CG solver




Performance

Some benchmarks - CG solver



Performance

Some benchmarks - CG solver




Performance

Some benchmarks - Matrix-Matrix Multiplication

Auto-tuning environment (AMD Radeon HD 7970, single precision)

GFLOP Performance for GEMM (Higher is Better)



What have we learned?

- What are subvectors/submatrices and how to use them

- How to eliminate temporaries

- Expression templates and when they help us

- Interface to Eigen

- ViennaCL isn't optimized for small vectors/matrices

- Performance bottleneck

- Overview of ViennaCL performance



ViennaCL: Behind the curtain



What to expect

- Backends

- OpenCL kernel management

- Extending ViennaCL

- ViennaCL and OpenGL

- Summary



Backends

There is more than OpenCL

CUDA from NVIDIA

OpenACC

Each framework has advantages and disadvantages



OpenCL

```
const char *kernel_string =
"__kernel void mykernel(__global double *buffer) {
    buffer[get_global_id(0)] = 42.0;
};";

int main() {
    ...
    cl_program my_prog = clCreateProgramWithSource(
        my_context, 1, &kernel_string, &source_len, &err);
    clBuildProgram(my_prog, 0, NULL, NULL, NULL, NULL);
    cl_kernel my_kernel = clCreateKernel(my_prog,
        "mykernel", &err);
    clSetKernelArg(my_kernel, 0, sizeof(cl_mem), &buffer);
    clEnqueueNDRangeKernel(queue, my_kernel, 1, NULL,
        &global_size, &local_size, 0, NULL, NULL);
}
```

Additional boilerplate code required (low-level API)

Broad hardware support (separate SDKs)

No more development effort from NVIDIA



NVIDIA CUDA


```
// GPU kernel:
__global__ void kernel(double *buffer)
{
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    buffer[idx] = 42.0;
}

// host code:
int main()
{
    ...
    cudaMalloc((void**)&buffer, size);
    kernel<<<blocknum, blockDim>>>(buffer);
    ...
}
```

Almost no additional code required

Vendor-lock

Relies on `nvcc` being available




OpenACC

```
void func(...) {  
    #pragma acc data pcopyin(A[0:size][0:size])  
    {  
        #pragma acc kernels loop  
        for(int i=0; i< size; i++)  
            for(int j=0; j < size; j++)  
                A[i][j] = 42;  
    }  
}  
  
int main()  
{  
    double A[1337][1337];  
    func(A);  
}
```

Simple OpenMP-type pragma annotations

Compiler support?

Insufficient control over memory transfers?



What to use?

Why choose one when we can support all?

ViennaCL has a backend layer

Backend is responsible for hardware interaction

Not only OpenCL anymore

Since ViennaCL 1.4.0

Different backends supported

OpenCL

OpenMP

CUDA



Backend support has to be enabled explicitly

```
viennacl::vector<float> v1, v2;  
v1 += v2;
```


CPU used!



Backend support has to be enabled explicitly

```
#define VIENNACL_WITH_OPENCL  
  
viennacl::vector<float> v1, v2;  
v1 += v2;
```

Now we are using OpenCL



Backends

Lets take a look!



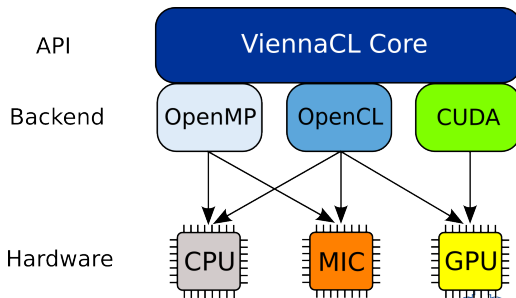
Vector Addition

Memory buffers can switch memory domain at runtime

```
void avbv(...) { // x = y + z
  switch (active_handle_id(x))
  {
    case MAIN_MEMORY:
      host_based::avbv(...);
      break;
    case OPENCL_MEMORY:
      opencl::avbv(...);
      break;
    case CUDA_MEMORY:
      cuda::avbv(...);
      break;
    default:
      raise_error();
  }
}
```


Memory Buffer Migration

```
vector<double> x = zero_vector<double>(42);  
  
memory_types src_memory_loc = memory_domain(x);  
switch_memory_domain(x, MAIN_MEMORY);  
/* do work on x in main memory here */  
switch_memory_domain(x, src_memory_loc);
```



Memory buffer switching at runtime

```
#define VIENNACL_WITH_OPENCL
#define VIENNACL_WITH_OPENMP

viennacl::vector<float> v1, v2;

switch_memory_domain(v1, MAIN_MEMORY);
switch_memory_domain(v2, MAIN_MEMORY);

v1 += v2; \\ working on CPU with OpenMP

switch_memory_domain(v1, OPENCL_MEMORY);
switch_memory_domain(v2, OPENCL_MEMORY);

v1 += v2; \\ working on GPU with OpenCL
```



OpenCL kernel management

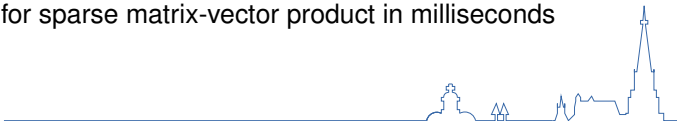
Best kernel implementations depend on target hardware

NVIDIA, AMD, Intel

Best work group size depends on target hardware

	NVIDIA				AMD			
	32	64	128	256	32	64	128	256
64	191	151	174	193	324	262	256	249
128	194	177	195	214	357	290	272	247
256	161	164	195	214	307	264	256	248
512	145	157	198	211	282	255	258	253

Execution times for sparse matrix-vector product in milliseconds



Kernel parameter tuning

Default number of work groups = 128

Default number of work items per work group = 128

Automatic tuning environment \Rightarrow XML file

How to use kernel parameters

```
using namespace viennacl;  
using viennacl::io;  
  
read_kernel_parameters< vector<float> >  
    ("float_vector_parameters.xml");  
read_kernel_parameters< matrix<float> >  
    ("float_matrix_parameters.xml");  
read_kernel_parameters< compressed_matrix<float> >  
    ("float_sparse_parameters.xml");
```



ViennaCL expression template don't cover all operations

Sample operation: $\mathbf{x} = \mathbf{A} \times [(\mathbf{y} \cdot (\mathbf{y} + \mathbf{z}))\mathbf{y} + \mathbf{z}]$

Automated kernel generation

Supported since ViennaCL 1.3.0

Experimental support

Symbolic variables

Operation is defined with C++ symbolic variables

Custom kernel object is generated



Automated kernel generation

Sample operation: $\mathbf{x} = \mathbf{A} \times [(\mathbf{y} \cdot (\mathbf{y} + \mathbf{z}))\mathbf{y} + \mathbf{z}]$


```
// Instantiation of the symbolic variables
symbolic_vector<NumericType, 0> sX;
symbolic_matrix<NumericType, 1> sA;
symbolic_vector<NumericType, 2> sY;
symbolic_vector<NumericType, 3> sZ;

//Creation of the custom operation
custom_operation my_op(
    sX = prod(sA, inner_prod(sY, sY+sZ) * sY + sZ),
    "operation_name" );
```



Automated kernel execution

```
viennacl::vector<NumericType> x, y, z;  
viennacl::matrix<NumericType> A;  
  
// fill data here  
  
//Execution of the custom operation  
viennacl::ocl::enqueue(my_op(x,A,y,z));
```



Not Everything Covered by ViennaCL

Complicated vector expressions in a single compute kernel

Direct OpenCL Kernel Handling is a Pain

```
const char * my_kernel_sources =
"__kernel void element_prod(\n"
"    __global const float * vec1,\n"
"    __global const float * vec2, \n"
"    __global float * result,\n"
"    unsigned int size) \n"
"{ \n"
"    for (unsigned int i = get_global_id(0); \n"
"        i < size; \n"
"        i += get_global_size(0))\n"
"        result[i] = vec1[i] * vec2[i];\n"
"};\n";
```


Extending ViennaCL

The OpenCL Way (error checks and casts omitted)

```
size_t source_len = std::string(my_compute_program).length();
cl_program my_prog = clCreateProgramWithSource(my_context, 1,
                                              &my_kernel_sources, &source_len, &err);
err = clBuildProgram(my_prog, 0, NULL, NULL, NULL, NULL);


const char * kernel_name = "element_prod";
cl_kernel my_kernel = clCreateKernel(my_prog, kernel_name, &err);

err = clSetKernelArg(my_kernel, 0, sizeof(cl_mem), &mem_vec1);
err = clSetKernelArg(my_kernel, 1, sizeof(cl_mem), &mem_vec2);
err = clSetKernelArg(my_kernel, 2, sizeof(cl_mem), &mem_result);
err = clSetKernelArg(my_kernel, 3, sizeof(unsigned int), &vsize);
err = clEnqueueNDRangeKernel(queues[0], my_kernel, 1, NULL,
                             &global_size, &local_size, 0, NULL, NULL);
```

Issues

Access `my_kernel` at some other location in an application?

What to do with `my_prog`?



Extending ViennaCL

The ViennaCL Way (namespaces omitted)

```
program & my_prog =  
    current_context().add_program(my_kernel_sources,  
                                   "my_program");  
kernel & my_kernel = my_prog.add_kernel("element_prod");  
enqueue(my_kernel(vec1, vec2, result, vec1.size()) );
```

At any other Location within the Application

```
kernel & my_kernel = get_kernel(  
    "my_program", "element_prod");  
viennacl::ocl::enqueue(  
    my_kernel(vec1, vec2, result, vec1.size()) );
```

Allows for Adding Missing Functionality Easily

A bit of OpenCL knowledge required



Extending ViennaCL

Integrate ViennaCL into User-Environment

User-provided context, queue and device

```
cl_context my_context = ...; //a context  
cl_device_id my_device = ...; //a device in my_context  
cl_command_queue my_queue = ...; //a queue for my_device  
// supply existing context 'my_context' with one device  
// and one queue to ViennaCL using id '0':  
viennacl::ocl::setup_context(0, my_context, my_device,  
    my_queue);
```

Wrapping Memory Buffers

```
cl_mem my_memory = ...;  
viennacl::vector<float> my_vec(my_memory, 10);
```

Use ViennaCL operations as usual



ViennaCL and OpenGL

Since OpenCL 1.1: OpenGL interoperability

With own OpenCL context: easy task

Workflow

Setup OpenGL and OpenCL

Create OpenGL buffer and OpenCL memory object

Pass OpenCL memory object to ViennaCL

Do ViennaCL magic

Use data in OpenGL



Setup OpenGL context (simple glut-glew magic)

```
glutInit(&argc, argv);  
  
glutInitDisplayMode(...);  
glutInitWindowPosition(100,100);  
glutInitWindowSize(1600,800);  
glutCreateWindow("CL - GL");  
  
glewInit();
```



Setup OpenCL context with OpenGL interoperability support

```
cl_context_properties properties[] = {
    CL_GL_CONTEXT_KHR, (cl_context_properties)
        glXGetCurrentContext(),
    CL_GLX_DISPLAY_KHR, (cl_context_properties)
        glXGetCurrentDisplay(),
    CL_CONTEXT_PLATFORM, (cl_context_properties)
        viennacl::ocl::get_platforms()[0].id(),
    0};

cl_device_id my_device =
    viennacl::ocl::current_device().id();

cl_context my_context = clCreateContext(properties, 1,
    &my_device, NULL, NULL, &err);
cl_command_queue my_queue = clCreateCommandQueue(
    my_context, my_device, 0, &err );
```



Setup OpenCL context with OpenGL interoperability support

```
cl_context_properties properties[] = {
    CL_GL_CONTEXT_KHR, (cl_context_properties)
        glXGetCurrentContext(),
    CL_GLX_DISPLAY_KHR, (cl_context_properties)
        glXGetCurrentDisplay(),
    CL_CONTEXT_PLATFORM, (cl_context_properties)
        viennacl::ocl::get_platforms()[0].id(),
    0};


cl_device_id my_device =
    viennacl::ocl::current_device().id();

cl_context my_context = clCreateContext(properties, 1,
    &my_device, NULL, NULL, &err);
cl_command_queue my_queue = clCreateCommandQueue(
    my_context, my_device, 0, &err );
```



Setting up ViennaCL for our context

```
viennacl::ocl::setup_context(  
    0,           // the ViennaCL context ID  
    my_context, // our context with OpenGL  
               // interoperability  
    my_device,  // the device we are working on  
    my_queue ); // the command queue for our context  
  
// tell ViennaCL that we want to use our context  
viennacl::ocl::switch_context( 0 );
```



Create OpenGL buffer and OpenCL memory object

```
glGenBuffers(1, &gl_buffer);  
glBindBuffer( GL_PIXEL_UNPACK_BUFFER, gl_buffer );  
glBufferData( GL_PIXEL_UNPACK_BUFFER, size, NULL,  
              GL_DYNAMIC_COPY );  
  
cl_mem cl_buffer = clCreateFromGLBuffer(my_context,  
                                         CL_MEM_READ_WRITE, gl_buffer, &err);
```



ViennaCL magic

```
// Create viennacl vector from OpenCL memory
viennacl::vector<float> my_vec(cl_buffer, size);

// Acquire memory object for write read/write operation
clEnqueueAcquireGLObjects(my_queue, 1, &cl_buffer,
    0, NULL, NULL);

// copy CPU data to ViennaCL
viennacl::copy( tmp_vec, my_vec );
// doing some stuff
my_vec *= 0.5f;

// Release memory object
clEnqueueReleaseGLObjects(my_queue, 1, &cl_buffer,
    0, NULL, NULL);
```



ViennaCL magic

```
// Create viennacl vector from OpenCL memory
viennacl::vector<float> my_vec(cl_buffer, size);


// Acquire memory object for write read/write operation
clEnqueueAcquireGLObjects(my_queue, 1, &cl_buffer,
    0, NULL, NULL);

// copy CPU data to ViennaCL
viennacl::copy( tmp_vec, my_vec );
// doing some stuff
my_vec *= 0.5f;

// Release memory object
clEnqueueReleaseGLObjects(my_queue, 1, &cl_buffer,
    0, NULL, NULL);
```



Lets take a look at this example



What have we learned?

- ViennaCL has different backends

- How to enable and use backends

- OpenCL management in ViennaCL

- How to use an own OpenCL kernel with ViennaCL

- How to provide own OpenCL contexts

- ViennaCL works with OpenGL!

- How to use ViennaCL to work with OpenGL



Acknowledgements

Contributors

Thomas Bertani

Evan Bollig

Philipp Grabenweger

Volodymyr Kysenko

Nikolay Lukash

Günther Mader

Vittorio Patriarca

Florian Rudolf

Astrid Rupp

Karl Rupp

Philippe Tillet

Markus Wagner

Josef Weinbub

Michael Wild



High-Level C++ Approach of ViennaCL

Convenience of single-threaded high-level libraries (Boost.uBLAS)

Header-only library for simple integration into existing code

MIT (X11) license

<http://viennacl.sourceforge.net/>

Selected Features

Backends: OpenMP, OpenCL, CUDA

Iterative Solvers: CG, BiCGStab, GMRES

Preconditioners: AMG, SPAI, ILU, Jacobi

BLAS: Levels 1-3

